

Text Detection between an AI Written Passage vs. a Human Written Passage

Saif Khan

Computer Science Department

Minot State University

Minot, ND, 58707

Saif.khan@minotstateu.edu

Kaif Khan

Computer Science Department

Minot State University

Minot, ND, 58707

Kaif.khan@minotstateu.edu

Prof. Muhammad Abusaqer

Computer Science Department

Minot State University

Minot, ND, 58707

Muhammad.abusaqer@minotstateu.edu

Abstract

In this research, we will dive into AIGT (AI-Generated Texts) detection with a specific focus on differentiating passages written by Artificial Intelligent (AI) models from those crafted by human writers. With the increase of advanced AI language models such as ChatGPT and Google's Bard, the need for robust methods to discern AI-generated text from human-authored content has become increasingly essential.

Our research uses a careful methodology to answer the question, what is the best way to train and detect AIGT from human-generated texts? We explored three different ways to evaluate this and figure out which one would give the most accurate results. (1) Classification methodology using Deep Learning (DL) with the help of BERT, and (2) Comparison between Machine Learning (ML) tools like Naive Bayes and Deep Learning tools like BERT to test which method can give the most accurate and consistent results. In addition, we tested the possibility of using Sentiment analysis as a tool to distinguish AIG and human writers. Leveraging a diverse dataset comprising texts generated by a range of

AI models and texts authored by human writers, our dataset includes 500,000 such essays. We reported Accuracy Precision, Recall, and F1 score, for our study experimentation.

The findings of this study have significant implications for applications such as plagiarism detection, content verification, and quality assurance in digital content creation in the future. By providing insights into the effectiveness of text detection methods, our research contributes to advancing text analysis techniques and informs the development of more reliable and efficient text authentication processes.

Keywords: Text detection, Artificial Intelligence, AI-generated content, Human-authored content, Comparative analysis, Content Authentication, Plagiarism Detection, Digital Content Analysis.

1 Introduction

Artificial Intelligence (AI) is a rapidly evolving field, and the advancements in Natural Language Processing (NLP) and Large Language Models (LLMs) have brought about text generation capabilities that very closely resemble human writing. Such technological advancement raises some serious red flags, particularly regarding the spread of plagiarism across various facets of modern-day life. Plagiarism, both in educational and professional environments, poses a serious threat to integrity and originality. The rise of ChatGPT from OpenAI, Gemini from Google, and Co-Pilot from Microsoft have made generating texts in large quantity incredibly accessible. In educational settings, the spread of AI-generated texts has made it challenging to distinguish between authentic student work and plagiarized content. This undermines the academic integrity of institutions but also robs genuine learners of the opportunity to showcase their own knowledge and creativity. Moreover, it erodes the foundation of scholarly discourse and intellectual advancement, as original contributions are overshadowed by the spread of copied content.

In our research we dive into the challenges of distinguishing between passages authored by humans and those generated by AI. Our objective is to determine the most optimal approach for constructing and training a model capable of reliably predicting texts generated by artificial intelligence (AIGT). Our aim is to develop a robust methodology that can ensure accurate identification of AIGT, thereby addressing the growing need for effective text authentication in contemporary contexts. Through the help of tools such as Naive Bayes, which is a powerful machine learning (ML) algorithm that is a probabilistic classifier, it evaluates the likelihood of certain words or features co-occurring to classify text effectively. Bidirectional Encoder Representations from Transformers (BERT), which is a sophisticated deep learning (DL) model developed by Google that captures contextual information in text, allowing for more precise analysis. We conducted a comparative analysis of both techniques to assess their relative effectiveness. Furthermore, we investigated using sentiment analysis as another way to figure out if a text was written by AI. Sentiment analysis involves examining the emotional tone behind the text, which can sometimes be different between human and AI-generated content. By studying sentiment patterns, we can improve the effectiveness of our detection methods. Overall, our goal is to contribute to the development of reliable methods for distinguishing between human and AI-generated passages. By testing these cutting-edge techniques, we aim to address the growing need for accurate text authentication models.

2 Literature Review

As the AI trend increases exponentially, detecting the difference between AI text and human text will become more demanding. It's necessary that we start training new models that can tell the difference between AI text and human text accurately. This research by Mindner, Schlippe, and Schaaf explores methods to identify text created entirely by AI and text that has been rephrased by AI from an original source. The authors employed a combination of features to train their detection systems, including perplexity, semantic analysis, readability scores, and feedback from other AI models. Their systems achieved high accuracy over 96% F1-score in classifying both basic and more sophisticated human-written and AI-generated texts. Additionally, the systems achieved an accuracy of over 78% F1-score for distinguishing between human-generated and AI-rephrased text. However, the authors acknowledged that as AI language models improve, detection will become more challenging, cautioning that the task of differentiating human and AI authorship is likely to become increasingly difficult as AI language capabilities advance. [1].

Another research done by Xiaomeng Hu, Pin-Yu Chen, Tsung-Yi Ho proposes RADAR, a framework that leverages adversarial learning to train an AI text detector, involving a unique approach: training a detector and a paraphraser in opposition to each other. The paraphraser learns to generate realistic text that evades detection as AI-generated, while the detector hones its ability to identify such text. This competitive training process enhances the robustness of the detector against paraphrased AI text. RADAR employs three neural networks: a target large language model, a detector, and a paraphraser. The detector is then trained to distinguish only AI-generated text, while the paraphraser attempts to create realistic text that bypasses detection. Through a feedback loop, the detector's evaluations inform the paraphraser's updates, and vice versa, strengthening both models iteratively. Hu, Chen, and Ho's (2023) experiments demonstrated that RADAR outperformed existing AI text detection models, particularly when dealing with paraphrased text. Extensive testing across various datasets and large language models proved RADAR's effectiveness in addressing the challenges associated with AI-generated content detection. However, it is important to acknowledge that the research is limited to some specific LLMs used for training and evaluation, the effectiveness of RADAR in the future may change [2].

Pengyu Wang introduced SeqXGPT, a novel method for fine-grained, sentence-level detection of AI-generated text, focusing on utilizing log probability lists from white-box large language models to identify patterns indicative of AI-generated text. SeqXGPT departs from previous methods by employing log probability lists, capturing the internal uncertainty of an LLM when generating text. The approach involves three key steps: Perplexity Extraction and Alignment, where perplexity lists are generated from a pre-trained LLM for each sentence Feature Encoding, utilizing a combination of convolutional and self-attention neural network layers to learn complex patterns distinguishing human-written and AI-generated sentences and Linear Classification, where a classification layer analyzes encoded features to label sentences as human-written or AI-generated. Wang demonstrated that his method achieved high precision, recall, and F1-score in detecting AI-

generated sentences. While excelling at identifying statistical inconsistencies within AI-generated text, SeqXGPT could benefit from integrating semantic information to enhance discernment of human-like AI-generated sentences. Despite this, the model presents a promising avenue for AI text detection, showcasing potential for addressing challenges associated with identifying AI-generated content at the sentence level while maintaining generalizability across domains [3].

Professor Chaka from University of South Africa examines the effectiveness of various AI content detection tools, including GPTZero and CopyLeaks, in identifying essays written by humans and AI, specifically focusing on the use of large language models like ChatGPT. The research emphasizes the importance of these tools in safeguarding academic integrity by combating plagiarism, particularly with the increasing sophistication of AI-generated essays. It employs a multi-faceted approach, evaluating the six AI text detectors on their ability to classify essays across four distinct writing styles: argumentative, descriptive, expository, and narrative. This approach provides valuable insights into the strengths and weaknesses of these detection tools across various writing formats commonly found in academic settings. The paper finds that all six detectors performed decently but not flawlessly in distinguishing human-written from AI-generated essays. While CopyLeaks exhibited the most promising results, none of the detectors achieved perfect accuracy. The author highlights limitations in current detection tools, due to their relatively new development stage. This finding demonstrates the need for continued research and development to enhance the accuracy of AI content detection tools as LLMs continue to evolve [4].

Research done by Junchao Wu shows the capabilities of large language models that have sparked a surge on methods to differentiate between human-written and AI-generated text, which is crucial for overcoming plagiarism and ensuring online content authenticity. Several key approaches that were employed by Wu are: Statistical Analysis involves analyzing properties like perplexity, which can indicate AI-generated content due to less variation in word choices compared to humans. Feature-Based Detection examines readability scores and word categories, with human-written text typically demonstrating broader vocabulary and sentence complexity. Adversarial Training pits a detector against a paraphraser, improving robustness against paraphrased text by enhancing the detector's ability to discern cleverly rephrased AI content. Inconsistency Detection targets illogical elements or factual inaccuracies, which may indicate AI authorship due to the difficulty in replicating human nuance and accuracy. These methods are promising but still under development, with ongoing refinement necessary to keep pace with advancements in AI text generation. Wu acknowledged that there are limited high-quality datasets to train these models on to get accurate results [5].

We have conducted an indepth analysis to compare BERT and Naive Bayes model, including sentiment analysis for detecting AI-generated text versus human-generated text. We noticed that BERT's result was an accuracy of 97.2%, F1-score of 96.2% and precision of 97.2% compared to Naive Bayes model which was lower with an accuracy of 95%. This shows that BERT model did better than Naive Bayes, but on the other hand BERT took way longer than imagined getting the results on only 1% of the entire dataset. Alternatively Naive Bayes was significantly faster running through all the 487,235 datasets. We also

added sentimental analysis that enhanced our research in figuring out and detecting AI text from human generated text. We go more detail in section 3.4.

3 Methodology and Results

In this section we start by explaining how we collected and explored our data using pandas. Then, in section 3.1 we dive into how we trained and tested the BERT model. In section 3.2 we discuss how we trained and built our Naive Bayes model. In 3.3 we dive into the comparative analysis of BERT and Naive Bayes. In 3.4 we look at how Sentiment Analysis can be used to detect AIGT.

Dataset Exploration

The dataset we used was sourced from Kaggle and provided by Shayan Gerami [6], which comprises of two primary columns: 'text' and 'generated'. The 'text' column contains randomly selected passages, each paired with a corresponding label in the 'generated' column. These labels indicate whether the passage was authored by a human ('0') or generated by an artificial intelligence ('1'). The dataset consists of 487,235 rows, with no NULL values in either column. Among these rows, 305,797 passages are labeled as human-authored, while 181,438 are attributed to AI generation. This distribution results in a ratio of 37.2% AI-generated passages to 62.7% human-authored passages within the dataset.

3.1 BERT Classification

In our research, we trained the powerful Bidirectional Encoder Representations from Transformers (BERT) model. The implementation of the BERT model involved utilizing several Python libraries, including: 1) Scikit-learn (sklearn) which offers a wide range of tools for data preprocessing, model selection, and evaluation. We chose sklearn for its efficient implementation of classification algorithms and metrics computation. 2) Pandas which is very commonly used for data manipulation. 3) PyTorch (torch) which is a deep learning framework, we used torch to build and train our BERT model. 4) From the Transformers library, we imported essential components for BERT model implementation, including BertForSequenceClassification and BertTokenizer. These modules provided pre-trained BERT models and tokenization functionalities necessary for sequence classification tasks. 5) AdamW optimizer and `get_linear_schedule_with_warmup` are key components for fine-tuning BERT models. AdamW offers efficient optimization with weight decay, while the scheduler adjusts the learning rate during training to improve convergence. Additionally, we employed performance evaluation metrics such as precision, recall, F1 score, and accuracy to assess the BERT model's effectiveness in classifying text sequences. Due to computational constraints we were facing, we opted to train the BERT model on a subset of the original dataset. We randomly sampled 1% of the dataset, resulting in 5000 data points. This down sampling expedited our training process, addressing the computational overhead associated with processing the entire dataset. During training, we set the number of epochs to 5, iterating over the dataset five times to optimize the model parameters. Despite the reduced dataset size, the training process took approximately 1 hour and 50 minutes to complete. Upon completion of training, we

evaluated the performance of the BERT model on a validation set. The model achieved impressive results, with a validation accuracy of 0.972, precision of 0.9729, recall of 0.952, and F1 score of 0.962.

3.2 Naive Bayes Classification

In our research, we utilized the Naive Bayes algorithm for text classification. Naive Bayes is commonly used in text classification tasks due to its computational efficiency, particularly with high-dimensional feature spaces found in NLP applications. Despite its simplistic assumptions, Naive Bayes often yields competitive performance and can provide quick and interpretable results. We implemented the Naive Bayes algorithm using Python libraries such as scikit-learn (sklearn) and pandas. We also used the TfidfVectorizer, short for Term Frequency-Inverse Document Frequency Vectorizer, which is a feature extraction technique that converts text documents into numerical vectors based on the frequency of terms and their importance across the corpus. It calculates a weight for each term in the document, considering both its frequency in the document and its rarity across the entire corpus. Also using Multinomial Naive Bayes (MultinomialNB) which is a variant of the Naive Bayes algorithm suitable for classification tasks with discrete features, such as word counts in text classification. It models the likelihood of observing each feature given the class and uses the probabilities to predict the most likely class label for a given sample.

Training and Evaluation

We trained the Naive Bayes model on the full dataset. The training process took approximately 63 seconds to complete, demonstrating the algorithm's efficiency in processing large datasets compared to more complex models like BERT. Upon training completion, we evaluated the model's performance using standard evaluation metrics such as accuracy, F1 score, and precision. The Naive Bayes model achieved an overall accuracy of 0.95, with an F1 score of 0.96 for human-authored passages and 0.93 for AI-generated passages. Additionally, the precision for human-authored passages was 0.93, while for AI-generated passages, it was 0.98.

3.3 Comparative Analysis between BERT and Naive Bayes

In conducting a comparative analysis between BERT and Naive Bayes models for distinguishing AI-generated from human-generated text, it's evident that BERT exhibits accuracy, with an impressive accuracy rate of 97.2%, along with a F1-score of 96.2% and precision of 97.2%, compared to Naive Bayes' accuracy of 95%. However, this advantage of BERT is that it gives higher accuracy because of its contextual understanding the algorithm is well written but can be mistaken by its time requirements depending on the size of the datasets, notably as it processed results on just 1% of the dataset, while Naive Bayes efficiently managed the entire dataset of 487,235 passages. Both models integrated sentiment analysis, improving their capacity to differentiate between AI-generated and

human-generated text, with BERT more likely leveraging its contextual comprehension for enhanced sentiment analysis. We believe that BERT may offer higher accuracy while taking more power, time and resources. At the same time, Naive Bayes is simple but is not nearly as accurate. In the real world we would use BERT because with the advancements of GPUs and faster computing power the model would be able to give off sharp results with faster timing.

3.4 Sentiment Analysis

In addition to machine learning techniques for text classification, we explored the utility of sentiment analysis as a supplementary tool for distinguishing between AI-generated and human-authored texts. Sentiment analysis, also known as opinion mining, aims to determine the sentiment or emotional tone conveyed in a piece of text, whether it be positive, negative, or neutral. By analyzing the underlying sentiments expressed in textual content, we sought to uncover potential indicators that could differentiate between AI-generated and human-authored texts. To conduct sentiment analysis, we employed the VADER (Valence Aware Dictionary and sEntiment Reasoner) sentiment analysis tool, a lexicon and rule-based sentiment analysis tool specifically designed for social media text. VADER is well-suited for analyzing short, informal text, making it a suitable choice for our investigation into AI-generated text detection. The tool provides sentiment scores for text passages, indicating the intensity of positive, negative, and neutral sentiments expressed within the text. Sentiment analysis has its limitations, particularly in contexts where textual content may exhibit ambiguity or sarcasm. Future research endeavors could explore more sophisticated sentiment analysis techniques, including deep learning-based approaches, to enhance the accuracy and robustness of sentiment analysis in detecting AI-generated texts. Additionally, integrating sentiment analysis with other text analysis methods, such as topic modeling and linguistic analysis, could further refine AI text detection models and improve their performance across diverse text genres and domains.

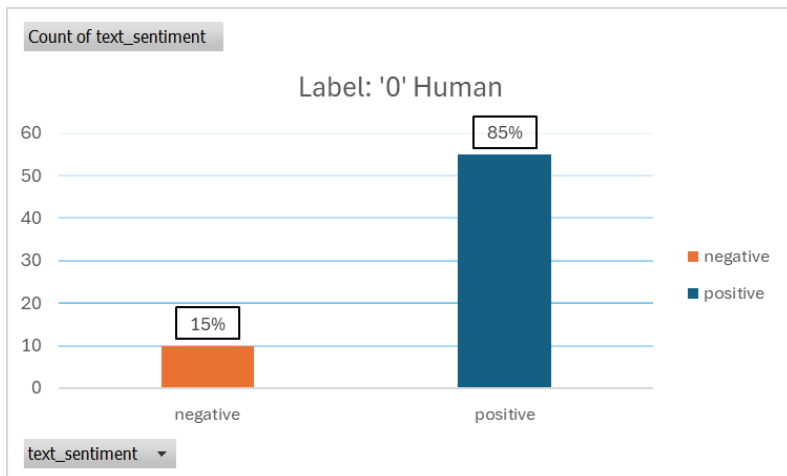


Figure 1: Sentiment Analysis on Human written passages. Indicates a higher negative sentiment.




Figure 2: Sentiment Analysis on AI written passages. Indicates a lower negative sentiment.

4 Conclusion

In this study, we delved into the realm of AI-generated text detection, leveraging machine learning techniques to discern between passages crafted by artificial intelligence models and those authored by humans. Our investigation centered on the comparative analysis of Naive Bayes and BERT, shedding light on their respective strengths and limitations in addressing the evolving challenges posed by advanced AI language models. Our research underscores the importance of robust text analysis techniques in the face of spreading AI-generated content. With the rise of sophisticated AI language models such as ChatGPT and Google's Bard, the need for reliable methods to differentiate between AI-generated and human-authored texts has become increasingly imperative. By evaluating the efficacy of Naive Bayes and BERT in this context, we contribute to advancing text analysis techniques and inform the development of more reliable and efficient text authentication processes.

5 Limitations

Despite the strides made in our research, several limitations warrant acknowledgment. The reliance on a single dataset and binary classification task may limit the generalizability of our findings across diverse text genres and AI models. Future research endeavors could explore more comprehensive datasets encompassing a wider range of text sources and writing styles. Additionally, investigating ensemble methods or hybrid approaches that integrate multiple classification techniques could further enhance the robustness and accuracy of AI-generated text detection models.

6 References

- [1] Mindner, Lorenz, et al. “Classification of Human- and Ai-Generated Texts: Investigating Features for Chatgpt.” *SpringerLink*, Springer Nature Singapore, 1 Jan. 1970, link.springer.com/chapter/10.1007/978-981-99-7947-9_12.
- [2] Hu, Xiaomeng, et al. “Radar: Robust AI-Text Detection via Adversarial Learning.” *arXiv.Org*, 24 Oct. 2023, arxiv.org/abs/2307.03838.
- [3] Wang, Pengyu, et al. “SEQXGPT: Sentence-Level AI-Generated Text Detection.” *arXiv.Org*, 15 Dec. 2023, arxiv.org/abs/2310.08903.
- [4] (PDF) *Detecting AI Content in Responses Generated by CHATGPT, ...*, www.researchgate.net/publication/372559876_Detecting_AI_content_in_responses_generated_by_ChatGPT_YouChat_and_Chatsonic_The_case_of_five_AI_content_detection_tools. Accessed 25 Mar. 2024.
- [5] Wu, Junchao, et al. “A Survey on LLM-Generated Text Detection: Necessity, Methods, and Future Directions.” *arXiv.Org*, 24 Oct. 2023, arxiv.org/abs/2310.14724.
- [6] Gerami, Shayan. “Ai vs Human Text.” *Kaggle*, 10 Jan. 2024, www.kaggle.com/datasets/shanegerami/ai-vs-human-text?resource=download.